

# SQL DML Cheat Sheet: SQLite | PostgreSQL | MySQL | SQL Server

## 1. INSERT

DB	Basic Syntax	Multi-row / SELECT Variant	Upsert / Conflict Handling	Notes & Gotchas
<b>SQLite</b>	INSERT INTO table (col1, col2, ...) VALUES (v1, v2, ...);	INSERT INTO table (cols...) SELECT ... FROM other_table;	INSERT OR REPLACE INTO ... or since SQLite 3.24+: INSERT ... ON CONFLICT ... DO UPDATE	REPLACE deletes then reinserts (may reset ROWID, trigger cascades). Be careful with side-effects and foreign keys.
<b>PostgreSQL</b>	INSERT INTO table (cols...) VALUES (...);	INSERT INTO table (cols...) SELECT ...;	INSERT ... ON CONFLICT (unique_col) DO UPDATE ... or DO NOTHING	Powerful RETURNING clause to fetch inserted rows.
<b>MySQL / MariaDB</b>	INSERT INTO table (cols...) VALUES (...);	INSERT INTO table (cols...) SELECT ...;	INSERT ... ON DUPLICATE KEY UPDATE ... or REPLACE INTO ...	ON DUPLICATE KEY works for unique/PK conflicts. REPLACE is delete + insert (watch foreign keys).
<b>SQL Server (T-SQL)</b>	INSERT INTO table (cols...) VALUES (...);	INSERT INTO table (cols...) SELECT ...;	Use MERGE (WHEN MATCHED / WHEN NOT MATCHED) for “upsert” logic	Use OUTPUT to return inserted rows. Be aware of known quirks/bugs in MERGE.

**Tips for INSERT in general:** - Always specify the column list explicitly; don't rely on implicit ordering. - Omitting columns means you must supply

*all* columns in table order. - For large inserts, use batched inserts to avoid memory/log issues. - Be mindful of default values, computed columns, identity/autoincrement semantics. - In migrations, test differences in defaults, nullability, unique constraints, trigger behavior across DBs.

## 2. UPDATE

DB	Basic Syntax	Update with JOINS / References to other tables	Returning / Output	Notes & Gotchas
<b>SQLite</b>	UPDATE table SET col1 = expr1, col2 = expr2 WHERE condition;	Doesn't support UPDATE ... JOIN syntax; use correlated subqueries or WHERE EXISTS	Newer SQLite versions support RETURNING (3.35)	Omitting <b>WHERE</b> updates all rows. Be cautious when referencing same-table subqueries.
<b>PostgreSQL</b>	UPDATE table SET col1 = expr1, col2 = expr2 WHERE condition;	Supports UPDATE ... FROM other_table ... WHERE ... join syntax	Supports RETURNING * for post-update results	Under MVCC, update = create new tuple + mark old as dead. Watch out for rule system rewrites.
<b>MySQL / MariaDB</b>	UPDATE table SET col1 = expr1, col2 = expr2 WHERE condition;	Supports UPDATE t1 JOIN t2 ON ... SET t1.col1 = t2.col2 WHERE ...	No built-in RETURNING in older versions; limited newer support	Be cautious with ambiguous joins, order of evaluation, and missing <b>WHERE</b> clauses.

DB	Basic Syntax	Update with JOINS / References to other tables	Returning / Output	Notes & Gotchas
<b>SQL Server (T-SQL)</b>	UPDATE t SET col1 = expr1, col2 = expr2 FROM table t JOIN other_table o ON ... WHERE ...	Yes: UPDATE ... FROM join syntax is standard in T-SQL	Use OUTPUT inserted.*, deleted.*	OUTPUT is powerful but has restrictions (e.g. can't always be used with views). Be careful about cascading updates, triggers, and identity columns.

**General Update Gotchas:** - Always include **WHERE** (or join filter) to avoid updating all rows. - Be aware of side-effects via triggers or cascading foreign key **ON UPDATE**. - In expressions like **SET col = col + 1**, understand how DB snapshots values (most DBs compute RHS from original value, not sequentially). - Null semantics: **col = NULL** is never true; use **IS NULL**. - Concurrent updates: consider locking, isolation levels, “lost updates”. - Updating unique or primary key columns can break constraints.

### 3. DELETE

DB	Basic Syntax	Delete with JOINS / Multi-table support	Returning / Output	Notes & Gotchas
<b>SQLite</b>	DELETE FROM table WHERE condition;	Doesn't support DELETE ... JOIN syntax; use WHERE EXISTS or subqueries	Newer SQLite versions support RETURNING (3.35)	Omitting <b>WHERE</b> deletes all rows. Foreign key cascades may delete related rows.

DB	Basic Syntax	Delete with JOINS / Multi-table support	Returning / Output	Notes & Gotchas
<b>PostgreSQL</b>	DELETE FROM table WHERE condition;	Supports DELETE ... USING other_table syntax: DELETE FROM t1 USING t2 WHERE t1.col = t2.col AND ...	RETURNING * allowed	On large tables, many dead tuples may accumulate; VACUUM may be needed.
<b>MySQL / MariaDB</b>	DELETE FROM table WHERE condition;	Supports DELETE t1 FROM t1 JOIN t2 ON ... WHERE ...	No built-in return before newer versions	Deleting many rows in one transaction can bloat logs; use chunked deletes.
<b>SQL Server (T-SQL)</b>	DELETE FROM table WHERE condition;	Supports DELETE t1 FROM t1 JOIN t2 ON ... WHERE ...	Use OUTPUT deleted.*	Watch cascading deletes, triggers, lock escalation. TRUNCATE TABLE is different (no row-by-row, faster, bypasses triggers, cannot be rolled back in some contexts).

**General Delete Gotchas:** - Always use **WHERE** unless deleting *all* rows intentionally. - Cascading deletes from foreign keys may trigger large deletion chains. - Deleting a large number of rows in one go may hit transaction log limits; break into smaller batches. - In MVCC engines, delete leaves “dead” row

versions/tuples; cleanup (vacuum / compaction) is necessary. - Use `RETURNING` / `OUTPUT` to see exactly what rows were deleted, but consider performance and side effect costs.

---

## 4. Summary & Best Practices

- Basic DML (`INSERT` / `UPDATE` / `DELETE`) is broadly portable, but **extensions** (joins in updates, returning/outputs, upsert/merge) vary significantly across engines.
  - To write **portable** SQL, stick close to the common subset: `INSERT` with column lists, `UPDATE` & `DELETE` with simple `WHERE`, avoid engine-specific features unless wrapped behind abstraction.
  - Use `RETURNING` / `OUTPUT` to avoid extra round trips when supported — but guard such usage behind dialect checks.
  - Be cautious with `MERGE` / `UPSERT` / `REPLACE` statements — they often carry subtle behavior (side effects, ordering, constraint resolution).
  - Before executing destructive or bulk operations, always test your `WHERE` filters via `SELECT`.
  - For heavy workloads, batch updates/deletes, monitor transaction log usage, handles locking contention, and manage version cleanup (in MVCC systems).
  - Be aware of hidden effects: triggers, cascades, implicit commits, side-effects. Always understand the full chain of dependencies.
  - Treat DML in migrations with care. Build a comprehensive test suite to validate behavior across DB platforms, nulls, constraint violations, edge boundary cases.
-